

Succeeding as a Zoox Compute Test Engineer

The First Year — Judgment, Relationships, and Impact

May 2026

Contents

1	How to Use This Guide	2
2	The First 90 Days	3
2.1	Track 1 — The Technical Ramp (what to do with the hardware)	3
2.2	Track 2 — The Relational and Ownership Ramp (how trust gets built)	3
3	The Org Map: Who You Work With and How	5
4	Ramping Without Breaking the Line	6
5	Thinking Like a Senior — Your Fast Track to the Title	7
6	A Systematic Debugging Method	8
7	Working With Contract Manufacturers	9
8	Communication and Influence	10
9	Continuous Improvement as a Discipline	11
10	What the JD Implies But Doesn't Say	12
10.1	Lab instrumentation & physical-layer measurement (<i>hidden in “control test instruments”</i>) . . .	12
10.2	Software engineering at scale (<i>hidden in “develop scripts and code”, “build and release”</i>)	12
10.3	Data systems, statistics & continuous improvement (<i>hidden in “analyze results”, “yield”</i>) . . .	12
10.4	Working across all four phases & influencing the design (<i>hidden in “all phases”, “validate new products”</i>)	12
10.5	Reliability/stress integration (<i>implied by “all phases” + safety-critical compute</i>)	13
10.6	Safety-criticality mindset & traceability (<i>inherent to a robotaxi role; not in the JD</i>)	13
10.7	Behavioral / seniority expectations (<i>hidden in “work independently, manage priorities”</i>)	13
11	Landing the Tools You Bring	14
12	Pitfalls to Avoid	15
13	Living the Leadership Principles (For a Test Engineer)	16
14	A Learning Plan for the Platform	17
15	What “Good” Looks Like	18

1 How to Use This Guide

The Study Guide ([zoox-study-guide.pdf](#)) makes you technically dangerous on day one — the registers, the math, the interfaces. This guide is about everything *around* that technical work that actually determines whether you succeed: how you ramp (technically *and* relationally — both tracks live here), who you work with, how you debug under pressure, how you support contract manufacturers, how you communicate findings, how you drive improvement, and how you land the tools you bring without overstepping.

You're coming in as a Test Engineer, but you bring senior-level experience (~13 years) to the role. You interviewed for the senior title; the team chose to start you as a Test Engineer — which is a gift, not a demotion: less title pressure, more room to ramp, learn the line, and build trust before the spotlight is on you. With your depth, the track to senior should be *shorter than usual* — so treat this guide as your fast track to earning it. That experience already changes the job in a specific way: nobody is going to hand you tickets. Success is measured by judgment and ownership — can you take an ambiguous problem (“this new board needs test coverage”) and turn it into a deployed, trusted, fast test program with the right people bought in. The technical skill is table stakes; the senior-level judgment in this guide is what earns the promotion fast.

The one sentence for this guide: *Earn trust fast by shipping small correct things, then use that trust to make the test process measurably better — without breaking the line or the relationships while you do it.*

A grounding note for your situation. You've done this transition before, just at smaller scale: you walked into 2G with no controls background and built production servo systems; you built a station dashboard nobody asked for and got it adopted. Those instincts — self-direction, build-to-learn, solve-the-real-problem — are exactly right here. The new variables are scale (volume manufacturing, multiple sites), stakes (a robotaxi's compute is safety-critical), and interfaces (more teams, contract manufacturers). This guide is about scaling what you already do well.

2 The First 90 Days

Your first 90 days run on two tracks at once, and they’re complementary, not sequential. The technical ramp is what to learn and do with the *hardware* — get a station, build the platform map, own an interface, debug to root cause. The relational/ownership ramp is the *people-and-judgment* side — earn trust, map the org, change nothing big until you’ve earned the context, then demonstrate independent judgment. The same calendar drives both: the technical track gives you something concrete to ship; the relational track makes sure shipping it builds trust instead of friction. Read them together — each week below has a technical job *and* a people job.

The deep technical detail behind every hardware task below — the registers, the bring-up order, the BERT math, the per-lane margining — lives in the Study Guide; this is the *sequence* and the *judgment* for applying it on a live line.

2.1 Track 1 — The Technical Ramp (what to do with the hardware)

Week 1 — orient and observe.

- Get a station and the current Linux test image; get accounts/access (results DB, dashboard, source control, the issue tracker).
- Run the *existing* test program on a known-good unit end to end. Read its code. Map it to the four manufacturing phases (PCBA / module / system / vehicle). Find where results land.
- Build the platform’s topology in your head: pull `lspci -t`, `nvme list`, and `nvidia-smi topo -m` on a real unit; sketch the root ports, switches, retimers, GPUs, NVMe drives, GMSL deserializers, NICs, and CAN. *Don’t change anything* — read-only, learn the normal so you’ll recognize the abnormal.
- Run the toolkit you brought, in mock mode, and show one person. (See “Landing the Tools You Bring” below for how to introduce it without overstepping.)

Month 1 — own a corner.

- Take ownership of one interface’s test (likely PCIe — your strength). Understand its current coverage, limits, and failure history from the data.
- Reproduce a known failure mode deliberately and watch every counter, so you *trust* the test before you change it.
- Make one small, safe improvement: better failure logging, an AER decode on a fail, a clearer operator message. Land it through the team’s release process. First contribution = trust deposited.

Months 2–3 — improve and extend.

- Pick a real continuous-improvement target from the data: a top Pareto failure mode, a slow test step, or a coverage gap (e.g., “we pass on link-up but don’t measure per-lane margin”). Propose it with data, build it, correlate on golden units, release it.
- Support a new-board bring-up or a new HW generation with EE — the JD’s “support test and validation of prototype designs.” This is where the Study Guide’s PCIe bring-up checklist earns its keep.
- Be the person who, when a board fails intermittently under thermal load, reaches for the AER decode *and* the scope and closes it to a named cause.

What “ramped” looks like at 90 days (technical): you can take a new board from EE, stand up coverage across the right phases, set data-driven limits, deploy it to a line or a CM, and debug a hardware failure to a physical root cause — independently. That’s the JD, delivered.

2.2 Track 2 — The Relational and Ownership Ramp (how trust gets built)

The technical track above is necessary but not sufficient: *how* you ramp determines whether each shipped thing earns trust or burns it. The same 30/60/90 windows, viewed through people and judgment.

2.2.1 Days 1–30: Earn trust, learn the system, change nothing big

Your only goals this month are to become useful at the existing process and to map the human and technical system. Resist the urge to redesign anything yet — you don’t have the context, and “the experienced new hire who immediately wanted to rewrite everything” is a reputation that takes a year to undo (and the surest way to *slow* your track to senior).

- Learn the existing test program cold before proposing changes to it. Run it, read it, trace a unit through it, find its failure history in the data. Assume it’s the way it is for reasons you don’t see yet — then ask *why*, genuinely curious, not leading.
- Meet your interfaces (the org map below): your manager, the senior test engineers, the EE and SW/FW leads for compute, the NPI/manufacturing folks, quality. For each: what do they need from test, what frustrates them about test today, how do they like to communicate.
- Find the data and the dashboard. Where do results live, what does yield look like, what are the top failure modes right now. You think in data; start swimming in theirs.
- Ship one tiny thing. A clearer error message, a fixed flaky test, a small log improvement — landed through the real release process. The point isn’t the impact; it’s proving you can move something through the system correctly and safely.

2.2.2 Days 31–60: Own a corner, contribute visibly

- Take ownership of one interface or test area (PCIe is the natural fit). Become the person people ask about it.
- Make a real, data-backed improvement — a coverage gap you can prove matters, or a runtime win. Propose it with numbers, build it, correlate it on golden units, release it.
- Start being the debugger. When a hard failure shows up in your area, take it, work it to root cause, and write up what you found clearly. Each closed bug is trust deposited.

2.2.3 Days 61–90: Demonstrate independent judgment

- Run a continuous-improvement project end to end — pick it from the data (top Pareto failure, slow step, missing coverage), drive it, measure the before/after, present the win.
- Support a prototype/new-HW bring-up with EE. This is the senior signal: you can stand up coverage for hardware that has no test yet.
- Have a point of view on where compute test should go (more margining-based limits, better CM correlation, faster soaks via confidence targets) — informed by 90 days of context, not day-one opinions.

Notice the CI project and the bring-up appear on *both* tracks — that’s not redundancy. The technical track says “stand up coverage for hardware with no test”; the relational track says “do it visibly, with the data, in a way that signals judgment.” Same work, two lenses; doing both is what makes a single act of work also an act of trust-building.

What success looks like at 90 days, where the two tracks meet: your manager hands you an ambiguous, important problem (“this new board needs coverage”) and doesn’t worry about it — because you’ve proven you can take it to a deployed, trusted, root-caused result without breaking the line. That’s the whole game.

3 The Org Map: Who You Work With and How

You can't be effective without knowing who owns what and what each group needs from you. Approximate map (titles vary):

Group	What they own	What they need from you	How to work with them
Your test team (mgr, senior test engineers)	The test programs, stations, framework	A teammate who ships and doesn't break the line	Learn their conventions; ship small first; ask before refactoring shared code
Electrical Engineering (EE) / HW design	Schematics, layout, the boards	Sharp failure evidence; bring-up support	Speak their language (rails, lanes, SI); hand them decoded data, not "it failed"
Software / Firmware / BSP	Linux image, drivers, firmware	Tests that work on their image; clear repro of driver/FW bugs	Pin versions; when test breaks after an image update, isolate test-vs-image cleanly
NPI / Manufacturing / Operations	Running tests on the line	Fast, robust, operator-proof tests; quick failure triage	They're your users; design for them; respond fast when the line is down
Quality / Reliability	Yield, RMAs, corrective action	Data, traceability, root-cause rigor	Give them parameters not just verdicts; help close the loop on field returns
Contract Manufacturers (CMs)	Building boards at volume	Released, documented, correlated test programs; remote support	Robustness + logs + correlation; treat them as partners (whole section below)
Program / Project Management	Schedule, deliverables	Honest status, early risk flags	No surprises; flag schedule risk the moment you see it, with options

The relationship that most defines a compute test engineer is with EE. When your test finds a failure, EE decides whether it's a real defect, a test problem, or a design issue. The quality of your evidence sets the speed and tone of that conversation. "PCIe failed" puts the burden on you to prove it. "Lane 7 shows 40 Replay-Timer correctable errors/min that appear only above 70 °C, here's the margining sweep and the rail scope" puts a defect in front of EE that's hard to wave off. Be the test engineer whose failures EE believes.

The relationship that will *test your discipline most* is with SW/FW, because of one recurring ambiguity: when a test starts failing after a BSP/driver/firmware update, is it a test regression or did the new image break the hardware path? You resolve this by pinning and logging every version (kernel, BSP, driver, FW slot, test-program version) with every result, so you can answer "what changed" in seconds instead of a day of finger-pointing. The engineer who can say "the test code is byte-identical to last week; only the driver moved from X to Y, and the failure tracks the driver" is the one who keeps that relationship collaborative instead of adversarial.

A note on cadence: these aren't one-time introductions. The healthy pattern is a standing light touch — you in EE's bring-up reviews, a recurring yield/Pareto sync with quality, a quick channel with the CM's lead engineer — so issues surface as small early signals, not as escalations. The PM relationship runs on one rule: no surprises. A schedule risk flagged three weeks out with two options is a planning input; the same risk surfaced the day the deliverable is due is a fire. Senior reads as *predictable*, and predictability is mostly about when you raise things, not whether you hit every date.

4 Ramping Without Breaking the Line

A robotaxi compute line is a live, safety-relevant production system. The fastest way to lose trust is to break it — a bad test release that false-fails good units halts a line and costs real money and credibility. The discipline:

- Read-only first. Everything you do in week one is observation. Read code, read data, watch tests run. Don't touch a register, a limit, or a config on a production station until you understand the blast radius.
- Sandbox before production. Have a dev station / spare unit. Reproduce, experiment, and validate there. Never debug a hypothesis on a production line.
- Respect the release process. There's a reason test programs are versioned and gated. Use it even when your change is "obviously fine." The one time you skip it is the time it isn't.
- Golden-unit gate. Before any release, confirm your change still passes known-good units and fails known-bad ones. A test that's stricter than you intended (false fails) or looser (escapes) shows up here, not on the line.
- Change one thing, measure, then the next. When debugging or tuning, single-variable discipline. The temptation under line-down pressure is to change five things; that's how you "fix" it without knowing why and have it return next week.
- Roll out staged, not all-at-once. When you can, prove a release on one station (or one shift, or a small unit count) and watch the yield before pushing it to every station and every CM. A limit that false-fails shows up as a yield dip on the canary station — caught on one line, not discovered simultaneously across three sites. This is the manufacturing version of a canary deploy.
- Watch the yield after every release. A release isn't "done" when it ships; it's done when the post-release yield and failure-Pareto look like you predicted. A silent yield drop right after your change is the change, until proven otherwise — so look, don't assume.
- Know how to roll back. Every release should be revertible, and rollback should be the *first* move when a release misbehaves on the line — revert, restore flow, then debug at the bench. Diagnosing on a stopped line wastes money you can save by reverting first. Knowing you can undo cleanly is what lets you move faster safely.

Your "restart-from-any-keyword" snapshot feature at 2G is the same instinct applied to operators: design for recovery, assume things will fail mid-process, make the safe path the easy path. Bring that instinct to how you release and operate, not just how you code.

5 Thinking Like a Senior — Your Fast Track to the Title

A few principles that separate senior test judgment from script-writing. These are the things to *say* and *live* — they’re what the manager, senior engineer, and director interviews were really probing for, what the job rewards, and the clearest way to earn the senior title faster than the calendar would. You don’t need the title to think this way; thinking this way is how you get it.

Ship only good units. The test exists to protect the vehicle and the brand from a marginal board. When you’re tempted to loosen a limit to recover yield, the question is “does this let a real defect through?” Yield matters, but never by hiding escapes — those come back $10,000\times$ more expensive, and in a robotaxi they’re a safety issue.

Coverage, runtime, yield is a triad you balance, not maximize. More coverage costs runtime and can cost yield (more chances to fail); faster runtime can cost coverage; looser limits buy yield but risk escapes. Every test decision moves these three. Seniority is making that trade *consciously and with data*, and being able to explain the trade you chose.

Data over opinion, always. “I think the limit should be 50 mV” loses to “the fleet distribution is 20 ± 5 mV, C_{pk} against a 50 mV limit is 2.0, here’s the histogram.” Capture parameters, build the distribution, set the limit from it. This is your superpower — you already think this way (the bilinear torque model, R^2 thresholds). Apply it to limits.

A failure is a question, not a verdict. When a unit fails, the job isn’t “mark it bad” — it’s “what does this tell us.” Is it a real defect (which kind, where), a marginal unit, a test problem, or a process drift? Good test engineers treat every fail as information about the *process*, not just the unit.

The cheapest defect is the one caught earliest. Internalize the $10\times$ curve (the Study Guide’s Manufacturing-Test chapter): cost to find a defect rises $\sim 10\times$ per phase it escapes to (PCBA 1x, module 10x, system 100x, vehicle 1000x, field 10,000x). Your instinct should always be “what’s the earliest phase that can catch this,” and a low-grade discomfort whenever a defect is found later than it could have been.

Test the thing, not the test. Beware tests that pass because they don’t actually exercise the failure mode (a “BERT” on an idle link, a thermal test that never gets hot, a link test that checks enumeration but not error counters under load). A test that can’t fail isn’t a test. Always ask: what defect would this catch, and have I proven it catches it?

6 A Systematic Debugging Method

The Study Guide gave you the *reflex* — failure → enumerate → dmesg → counters → isolate → measure → decode to a part. Here's the explicit *method* for when a board fails and people are waiting. It works because it's evidence-driven and avoids the two failure modes of debugging under pressure: guessing, and changing many things at once.

1. Reproduce and define “failed.” Get a precise, repeatable failure. “Sometimes flaky” is not yet a bug — pin down the exact symptom, rate, and conditions (temperature? load? specific unit? specific lane/port?). If you can't reproduce it, your first job is making it reproducible.
2. Read what the system already recorded. dmesg, the test logs, the counters (AER, ECC, EDAC, SMART). Most hardware failures already told you what happened; read before you poke.
3. Localize by layer and by swap. Use the layer the error names (the Study Guide's PCIe and GPU chapters — AER correctable/uncorrectable bits, GPU XID codes, the error tables) to pick physical vs protocol. Then isolate by swapping: known-good unit in the failing slot, failing unit in a known-good slot — does the failure follow the unit or stay with the station/slot? This one technique resolves a huge fraction of “is it the board or the fixture” questions.
4. Form one hypothesis and test it with one change. “It's thermal” → soak hot and watch the counter. “It's the rail” → scope it under load. Single variable. Predict what you'll see *before* you look; a surprise is more informative than a confirmation.
5. Measure at the physical layer when counters run out. Scope the rail/ripple/sequencing, margin the lane, check termination. The counters tell you *what*; the instruments tell you *why*.
6. Get to a root cause you can name and a fix you can verify. Not “reseated it and it went away” — *why* did reseating fix it (connector contact? specific lane?), and does the fix hold across thermal cycles and multiple units. “Fixed, cause unknown” means it'll be back.
7. Write it down so the next person doesn't re-walk it. A short, clear writeup with the evidence. This is how a debug becomes institutional knowledge and how you build a reputation as the person who actually closes things.

Bisect, don't crawl. When the fault could be anywhere along a chain — a range of FW/driver versions, a long signal path with retimers and connectors, a sequence of test steps — halve the search space each step instead of walking it linearly. *Versions*: a failure that appeared between two releases is a `git bisect` over the versions, not a line-by-line code read. *Topology*: on a cabled board-to-board link, lane margining at the retimer's receiver (the Study Guide's PCIe chapter) localizes a marginal eye to “before vs after the retimer” — board trace vs cable — in one measurement. *Process*: swap good/bad across slot, fixture, cable, and unit to find which variable carries the failure. One good bisection step is worth ten guesses.

Know when to stop and escalate — with evidence, not a shrug. Single-variable rigor is not the same as working a problem alone forever. If you've localized the failure to a layer you don't own (a driver bug, a silicon erratum, a design marginality) or you're past the point of diminishing returns, the senior move is to hand it off *with the evidence package already built* — decoded errors, the swap matrix, the conditions, the bisection result — so EE or SW picks it up at step 5, not step 1. Escalating a well-characterized problem is closing it, not giving up; escalating “it's broken” is the thing to avoid.

The trap to avoid under line-down pressure is shotgunning — changing several things to make it go away. It sometimes “works” and always leaves you without a root cause, so it comes back. Slower-but-single-variable beats fast-but-guessing almost every time.

7 Working With Contract Manufacturers

The JD calls this out explicitly — “*build and release test solutions for use on the manufacturing lines at Zoox and at contract manufacturing partners.*” For a senior test engineer this is one of the highest-leverage and least-glamorous parts of the job, and it’s where robustness and communication matter more than cleverness.

Why CMs exist: Zoox doesn’t build every board in-house at volume. Contract manufacturers (the EMS companies — think Flex, Jabil, and the like) have the lines, capacity, and cost structure to build at scale. You release your test programs *to them*, and they run your tests on units you may never physically touch.

What changes when the line is 2,000 miles away:

- Robustness is everything. A test that occasionally hangs is an annoyance at your bench and a line-stopping incident at a CM where you can’t walk over and fix it. Defensive coding, timeouts on every instrument/DUT call, clear failure states, and graceful recovery aren’t nice-to-haves — they’re the difference between a 10-minute issue and an overnight outage across a time zone.
- Logs are your eyes. You debug CM failures through logs and remote access. Every test must log enough — measured values, decoded errors, dmesg snippets, versions — that you can diagnose a failure from the record alone. The “attach the evidence to the failure” habit the Study Guide teaches is non-negotiable here.
- Documentation and training. The CM’s operators and engineers run your program. They need setup docs, a clear runbook, fixture instructions, and a triage guide for common fails. You’ll train remotely; assume nothing is obvious.
- Correlation and golden units. The same unit must produce the same result at Zoox and at the CM — that’s cross-site gauge reproducibility. You ship golden units (characterized known-good and known-bad) to the CM and correlate: do their stations agree with yours? A correlation gap means a fixture, calibration, or environment difference you must resolve before trusting their yield.
- Their yield is partly your responsibility. A test that false-fails at the CM costs them throughput and costs you credibility; an escape that ships from the CM is your coverage gap. You own the test even where you don’t own the line.

The release package you hand a CM is the artifact that defines this relationship: versioned test code + config, setup/runbook docs, fixture spec, acceptance criteria, golden-unit correlation data, and a failure-triage guide. Treat assembling it as a core deliverable, not an afterthought.

The on-call / line-down reflex. A CM line stop is a cross-time-zone incident — the line is down, money is burning, and you can’t walk over. The same things that make a test robust make the incident short: timeouts on every instrument/DUT call so nothing hangs forever, clear failure states so the operator’s screenshot tells you the layer, logs rich enough to diagnose from the record, and a triage guide that lets *their* engineer resolve the common 80% without waking you. The difference between a 10-minute issue and an overnight outage is almost entirely robustness + logs + documentation you built in advance, not heroics at 3 a.m. Build for the incident you won’t be awake for.

Practical reality: time zones, language differences, and you-can’t-see-it. Over-communicate in writing, make specs unambiguous, and build tests that explain themselves. The senior move is making the CM *successful and self-sufficient*, not dependent on a call to you for every failure.

8 Communication and Influence

Most of your impact flows through other people — EE fixing a design, SW fixing a driver, ops adopting a test, a CM running it right. How you communicate determines whether that happens.

Lead with evidence, structured for the reader. A failure report to EE: symptom → conditions → evidence (decoded counters, margining, scope) → what it points to → what you need from them. A yield update to your manager: the number, the trend, the top Pareto cause, what you’re doing about it. Tailor the altitude — EE wants the lane and the rail; your director wants the yield and the risk.

Make the data visible. Your FastAPI station dashboard at 2G is the template: you turned “operators running between 4 stations” into “one engineer watching a browser,” and idle-after-failure time dropped from 30+ min to under 5. That’s influence through tooling — you didn’t argue for better monitoring, you *built* it and let it sell itself. Look for the Zoox-scale version of that (fleet yield/runtime visibility, failure-mode dashboards) once you’ve earned the context to build it.

Influence without authority. As an experienced individual contributor you’ll often need EE or SW to do something and have no authority to make them. What works: bring data not opinions, frame it as *their* problem solved (a defect they’d want caught, a driver bug repro’d cleanly), and make the ask small and specific. The dashboard story is the pattern — solve the real problem, show it, let adoption follow.

Communicate the incident, not the panic. When a line is down — yours or a CM’s — the update that calms the room has four parts and fits in five lines: impact (which line, how many units blocked, since when), status (what you know and what you’re doing right now), ETA or next checkpoint (“rolling back now, flow restored in ~15 min” or “next update in 30 min”), and the ask (what you need and from whom). Send it early, send it on a clock, and update on the cadence you promised. Silence during an incident reads as “out of control” even when you’re making progress; a steady, structured drumbeat reads as “this is handled.”

Deliver bad news early and straight. A slipping schedule, a coverage gap you found, a unit that escaped — surface it the moment you’re confident, with the impact and your proposed options, not after it’s unrecoverable. Engineers and managers forgive a problem raised early with a plan; they don’t forgive being surprised by one you sat on. “Here’s the issue, here’s what I recommend, here’s the call I need from you” is how a senior raises a problem.

Disagree and commit. You’ll sometimes lose an argument about a limit or a coverage call. State your case once, clearly, with data; if the decision goes the other way, commit to it fully and revisit with new data later if you’re proven right. Re-litigating lost decisions burns the trust you need for the next one.

Write it down. Findings, root causes, decisions, runbooks. Written artifacts scale across sites and time zones and outlive the conversation. The test engineer whose investigations are documented becomes the source of truth; the one who keeps it in their head becomes a bottleneck.

9 Continuous Improvement as a Discipline

The JD makes this an explicit responsibility: “*Establish best practices and drive continuous improvement for improving test deployment, test runtime, and yield.*” This is where you go from “runs the tests” to “makes the test process better” — the senior mandate.

Pick projects from data, not vibes. The fleet results tell you where the pain is: the top Pareto failure mode (yield), the slowest test step (runtime), the test that deploys slowly (deployment), the coverage gap that’s letting escapes through. Let the data nominate the project; that also makes the win measurable.

Run it as a closed loop. Baseline the metric → form a hypothesis → make the change → correlate on golden units → deploy → measure the before/after → report. A CI project without a measured delta is just a change. The measurement is what makes it a contribution and what makes the next one easier to get approved.

Translate the delta into the language the business runs on. “Cut the soak from 300 s to 90 s” is an engineering result; “freed 3.5 min/unit, which at this line’s takt is roughly one extra unit per station-hour” is an impact your manager can take upward. Runtime maps to throughput and capacity; yield maps to scrap/rework cost and units shipped; deployment speed maps to time-to-ramp a new board or CM. You don’t need precise dollars, but framing the win in throughput/yield/ramp terms is what turns “a nice optimization” into “a result with my name on it” — and it’s the difference the promotion case is built from.

Don’t optimize a metric that doesn’t matter — or game one that does. Shaving runtime off a test step that isn’t the takt bottleneck moves nothing; find the *constraint* first (the slowest step gating throughput) and work that. And never “improve yield” by quietly loosening a limit — that’s gaming the number while raising the escape rate, the cardinal sin on safety-critical compute. A real CI win improves the metric *without* trading away coverage or correctness; if a change helps one leg of the coverage/runtime/yield triad by hurting another, say so explicitly and make the trade consciously.

Examples sized for your first year:

- *Runtime*: replace a fixed, padded soak with a confidence-target BERT (the Study Guide’s Math chapter) — run exactly long enough to prove a BER below $1e-12$ at 95% confidence, then stop. Often a large, safe time win with a statistical guarantee.
- *Yield*: a top failure mode is “PCIe link marginal” with no diagnostic — add lane margining so marginal-but-passing units are caught earlier and real defects come with evidence, *and* set the limit from the fleet distribution instead of a guess.
- *Deployment*: move station-specific settings from code into config so a new board revision or a new CM is a config change, not a code release and re-qualification.

Best practices as artifacts. “Establish best practices” means turning your judgment into things the team reuses: a bring-up checklist, a release/correlation procedure, a failure-triage guide, a standard result schema. Codify it so it survives you and scales to the next hire.

10 What the JD Implies But Doesn't Say

The job description is a paragraph; the *job* is much larger. A senior compute test engineer is expected to bring a lot the JD only gestures at. This chapter makes the implied requirements explicit — each notes the JD phrase it hides behind — so none of them surprises you. Treat it as a checklist of “things I’m quietly expected to be good at.”

10.1 Lab instrumentation & physical-layer measurement (*hidden in “control test instruments”*)

- Bench-instrument fluency and *what each proves on a compute board*: rail voltages under load, ripple/noise, power-up sequencing, inrush, PERST#/reset timing (the Study Guide’s Power and Manufacturing-Test chapters). Not “I’ve seen a scope” — “I scope the rail AC-coupled under the load profile that triggers the AER burst.”
- Instrument automation stack: VISA/SCPI over LAN/LXI/USB/GPIB, and a shared-instrument server for multi-station benches (your `equipment_rpc.py`).
- Measurement discipline: 4-wire/Kelvin for low-R and shunts, settling/averaging, fixed vs autorange, and calibration/NIST traceability — log instrument IDN + cal status with every measurement, because a number from an out-of-cal instrument is not a result.
- The instinct that “digital failures are often power/SI failures” — reaching for the scope + AER decode together to root-cause an intermittent link error.

10.2 Software engineering at scale (*hidden in “develop scripts and code”, “build and release”*)

- Not “can write Python,” but test-framework engineering: a config-driven harness, pytest fixtures/parametrize, structured result schemas, defensive coding with timeouts on every instrument/DUT call, graceful failure/recovery.
- C/C++ for tight loops where Python can’t keep up (the AER clear-and-count, the BERT inner loop) — the JD’s “C++ or C# beneficial.”
- Config-over-code so a new board revision or a new CM is a config change, not a code release + re-qualification.
- Version control, release engineering, rollback for *test programs as released artifacts*.

10.3 Data systems, statistics & continuous improvement (*hidden in “analyze results”, “yield”*)

- SPC, C_{pk}/P_{pk} , GR&R/MSA, FPY/RTY, guard-banding, data-driven limit setting — the statistical backbone of “improve yield” (the Study Guide’s Manufacturing-Test and Math chapters have the detail). This is the difference between “I set the limit at 50 mV” and “the fleet is 20 ± 5 mV, C_{pk} against 50 mV is 2.0, here’s the histogram.”
- A results database + dashboards at fleet scale; Pareto + RCA (5-whys/fishbone) to pick and close the top failure modes; closed-loop CI (baseline → change → correlate → deploy → measure delta).
- MES / test-data-management integration awareness (check-in/out, push verdicts + parameters).

10.4 Working across all four phases & influencing the design (*hidden in “all phases”, “validate new products”*)

- Understanding the coverage envelope of PCBA methods (AOI/AXI/ICT/JTAG/flying probe) even though the CM runs them, so you can allocate coverage across phases (the $10\times$ placement skill).
- DFT (Design-for-Test) influence on the design — pushing *upstream*, while EE is still laying out the board, for test pads/ICT access, boundary-scan/JTAG coverage, a scratch/ID register on custom FPGA cards, accessible rails, telemetry hooks.
- Bring-up of custom hardware with no vendor test plan — working from the schematic (root-port mapping, bifurcation, retimers, rails) to stand up coverage from scratch.

10.5 Reliability/stress integration (*implied by “all phases” + safety-critical compute*)

- Knowing where HALT/HASS/ESS/burn-in/thermal-cycling belong (HALT = design tool; HASS/ESS/burn-in = production screens) and that the test engineer’s job is the in-soak functional monitor — the at-temperature link/error/throttle checks during the screen.
- The “passes at 25 °C, fails at 85 °C” intuition (SerDes margin loss + Arrhenius) and therefore testing hot.

10.6 Safety-criticality mindset & traceability (*inherent to a robotaxi role; not in the JD*)

- “Ship only good units” as non-negotiable: never recover yield by raising the escape rate on safety-critical compute.
- ISO 26262 awareness: ASIL levels, mandatory traceability, coverage/limits as safety-case evidence, change control on test programs/limits.
- Full traceability/genealogy expectation: serial → genealogy → program version → parameters → disposition as an auditable chain.

10.7 Behavioral / seniority expectations (*hidden in “work independently, manage priorities”*)

- Judgment and ownership over an ambiguous mandate (“this board needs coverage”) rather than ticket-taking; the coverage/runtime/yield triad traded *consciously, with data*.
- Evidence-based communication (decoded, layer-identified failures EE believes; altitude-matched updates), influence without authority, disagree-and-commit, and writing it down so knowledge scales across sites.
- Operator/CM-centric design: operators and CM engineers are your *users*; an operator-hostile test fails in practice no matter how technically correct.
- Not breaking the line: read-only-first, sandbox-before-production, golden-unit gate, single-variable debugging, rollback.

How to use this list. These are the unspoken bar. You don’t recite them — you *demonstrate* them: a failure report that’s already decoded and layer-identified; a limit backed by a distribution and a C_{pk} ; a test that’s robust enough to run unattended at a CM; a bring-up done from a schematic with no vendor plan. Every one of them is “senior” rendered as a concrete behavior.

11 Landing the Tools You Bring

You're walking in with a real toolkit (the PCIe BERT, AER decode, lane margining, the NVMe/GPU/GMSL checks, the harness). Done right, it's a fast credibility win; done wrong, it's "the new person who ignored our stack." The sequence that works:

1. Learn their framework first. Before proposing your tool, understand what Zoxx already has. They may already have a BERT, a harness, a dashboard. Adopt their conventions; your goal is to *add to* their stack, not replace it.
2. Find the genuine gap. Your strongest adds are likely the confidence-target BERT (runtime + rigor), lane-margining-based limits (a coverage upgrade over pass-on-link-up), and the clean AER-decode-with-evidence habit. Lead with the gap their current process has, not with "I built a thing."
3. Demo in mock mode, framed as a question. "I built this at home to keep my PCIe skills sharp — it runs the BERT to a confidence target and does receiver lane margining. Could something like this help with X?" is collaborative. "Replace your tool with mine" is not. The mock backend means you can show it working on your laptop on day one with zero risk.
4. Get a senior engineer's buy-in early. Make an ally of whoever owns the current framework. Contribute your tool *into* their structure, with their review, on their release process. Shared ownership beats a parallel tool nobody else maintains.
5. Be gracious about "we already do that." Sometimes the value is one piece (the confidence math, the margining) inside their existing flow, not the whole tool. Take the win that's real. Contributing a useful piece to the team's tool earns more than insisting on your whole tool.

The toolkit's real day-one value isn't "deploy it to the line tomorrow" — it's proof of how you think (read-only-safe, mock-testable, evidence-attached, data-driven limits) and a concrete artifact for the "what would you build" conversations. That's worth a lot even if not one line ever runs on a Zoxx station.

12 Pitfalls to Avoid

The common ways a strong, experienced new hire stumbles in the first months — forewarned:

- Rewriting everything immediately. You don't have the context yet. Earn it first; the existing system has reasons. Ship small, then reshape.
- Breaking the line. A careless test release that false-fails good units is the fastest trust-destroyer there is. Golden-unit gate + release process + rollback, every time.
- "It failed" with no evidence. Forces others to reproduce and erodes your credibility. Always attach decoded, layer-identified evidence.
- Over-promising on schedule. "No surprises" beats "optimistic then late." Flag risk early with options. Senior = predictable.
- Shotgun debugging. Changing many things under pressure. Single variable, named root cause, verified fix — or it comes back.
- Chasing yield by loosening limits without checking escapes. Recovering yield by letting defects through is the cardinal sin in a safety-critical product.
- Treating operators/CMs as afterthoughts. They're your users. A test that's technically correct but operator-hostile fails in practice.
- Version confusion. "Test broke" after a firmware/image update — pin and log versions so you can instantly tell test-regression from HW/FW change.
- Hero culture. Keeping knowledge in your head makes you a bottleneck and a single point of failure. Write it down; make yourself replaceable, which paradoxically makes you promotable.
- Tool-first instead of problem-first. Bringing a solution looking for a problem. Find the real gap, then apply the tool.

13 Living the Leadership Principles (For a Test Engineer)

Zoox operates within Amazon, so the Leadership Principles are the cultural backbone and showed up in your interviews — and they’re not just interview theater: at Amazon/Zoox they are the literal axes a promotion case is written and evaluated against. So living them visibly is the mechanism by which a Test Engineer with senior-level depth gets recognized as senior. The point isn’t to recite them — it’s to *live* them in concrete test-engineering behavior, and to have a real story for each:

- Customer Obsession — your customers are the rider’s safety and the operator/CM who runs your test. “Ship only good units” is customer obsession for a robotaxi; an operator-proof, fast, clear test is customer obsession for the line.
- Ownership — own the test where you don’t own the line (CMs), own the failure to root cause, own the outcome not just the task. Think beyond “my script ran.”
- Dive Deep — root cause, not symptom; data not anecdote; read the registers and the distributions. This is your natural strength — make it visible.
- Insist on the Highest Standards — a test that can’t catch the defect isn’t done; a limit set by guess isn’t done; “fixed, cause unknown” isn’t done.
- Earn Trust — evidence-based findings, don’t break the line, do what you said. Trust is the currency that lets you drive change.
- Bias for Action — the dashboard you built unasked is the canonical example: see the problem, build the fix, show it. Balanced against not breaking the line (action $\neq$ recklessness).
- Invent and Simplify — the confidence-target BERT and margining-based limits are invention; config-over-code and a clean result schema are simplification.
- Are Right, A Lot / Have Backbone; Disagree and Commit — argue the limit with data, then commit to the decision; revisit with new data, don’t re-litigate. (Two distinct principles Amazon lists separately, but for a test engineer they’re the same muscle: be right because you’re data-driven, push when you believe it, commit once the call is made.)
- Frugality — you built the dashboard with no budget on a spare PC. Build-vs-buy judgment and doing more with less is a strength here, not a constraint; a confidence-target soak that frees station time is frugality with the line’s capacity.
- Learn and Be Curious — the whole 90-day ramp is this principle: read the framework, pull the topology, learn the silicon you haven’t touched (GMSL, the specific GPUs), ask EE *why* a limit is where it is. Your build-to-learn instinct (the toolkit you wrote at home) is this principle made visible.
- Deliver Results — the one the promotion case leans on hardest: not “I worked on PCIe test” but “I cut a module-test soak from 300 s to 90 s with a statistical guarantee and zero escape-rate increase, deployed across two sites.” Tie every effort to a measured outcome and you’re speaking the language seniority is judged in.

When a behavioral question comes (and they come woven into technical discussions, per the Study Guide), these are lived examples, not memorized values — which is exactly what they’re listening for. The strongest answers are STAR-shaped (Situation, Task, Action, Result) and land on a *measured* result, because that’s where Deliver Results and Dive Deep both get scored.

14 A Learning Plan for the Platform

You'll learn fastest by doing, but point the curiosity deliberately:

Read (in priority order): the existing test framework's docs and code; the compute platform's design docs and schematics (especially the PCIe topology — root ports, switches, retimers, bifurcation); the BSP/driver and firmware-update docs; the CM release/correlation process docs; and the relevant slices of the PCIe base spec (AER, link training, equalization, lane margining) and NVMe spec (admin commands, log pages, self-test) — you don't read these cover to cover, you read the chapters behind the registers you touch.

Learn the systems: the results database and its schema, the monitoring dashboard, source control and the release/deploy process, the issue tracker, and the lab — bench inventory, instrument calibration system, fixtures, golden units.

Learn from people: shadow a senior test engineer through a real debug; sit with EE during a board bring-up; watch operators run the test (you'll learn more about your test's real UX in an hour of watching than a week of reading); ask the quality team what field returns actually look like.

Build the platform map yourself: on a real unit, pull the full topology (`lspci -tv`, `nvme list`, `nvidia-smi topo -m`) and trace a unit through the four phases. The mental model of “what's connected to what and what's tested where” is the foundation everything else hangs on, and building it yourself cements it.

Produce artifacts as you learn, don't just absorb. You retain (and signal competence) by making things, not by reading. Concretely, in the first weeks aim to produce: a one-page topology diagram of the compute platform (root ports → switches/retimers → GPUs/NVMe/GMSL/ NICs); a failure-Pareto of your interface pulled from the results DB; a short “how the existing test program works” writeup mapping each step to a phase and a defect class; and a glossary of the local vocabulary (their names for stations, fixtures, board revs, the release process). These double as your learning record and as early, low-risk contributions the team can actually use.

A learning checkpoint at each window. By 30 days you should be able to draw the topology from memory and run the existing program unaided. By 60 days you should know your interface's coverage, limits, and top failure modes cold, and have read the spec chapters behind the registers you touch. By 90 days you should be able to stand up coverage for a *new* board from its schematic — the point where reading has turned into capability.

15 What “Good” Looks Like

Concrete bars to aim for, so you know if you’re on track:

At 90 days: You own an interface’s test (likely PCIe). You’ve closed real hardware bugs to named root causes. Your manager hands you ambiguous, important problems without worrying. You ship through the release process safely and haven’t broken a line.

At 6 months: You own a board’s or a generation’s test program end to end. You’ve driven at least one measured continuous-improvement win (runtime, yield, or deployment). You’ve correlated a test across a CM. You’re the person the team asks about PCIe and high-speed link failures.

At 1 year: You shape compute test *strategy*, not just execute it — coverage allocation across phases, margining-based limits, confidence-target soaks. You’ve brought a new capability to the team that’s now standard practice. You mentor newer engineers. There are measurable yield and runtime improvements with your name on them. EE and SW seek you out during bring-up because your evidence makes their job faster.

That arc — useful, then trusted, then shaping the strategy — is the whole job. You’ve walked a version of it before at smaller scale. Walk it again, bigger, and you’ll have done exactly the senior-level work you were brought in to do — and turned a Test Engineer title into a senior one faster than anyone expected.

Go get it.